

Python for Data Science

# Churn Prediction

Hae In Keum  
Julius Enderwitz



# Project overview

## Predicting User Churn from Behavioral Logs

- **Goal:** Predict user churn in a 10-day horizon using event-level log data
  - Transform the clickstream logs into user-level features without leakage → observe windows
- **Our approach:**
  - Window-based feature construction to capture recent behavior of individual users
  - Multiple backward-sliding snapshots per user to maximize training data usage
  - Use tree-based ensemble models (LGBM) to capture non-linearities and interactions
  - Explicit threshold tuning on the validation set to handle class imbalance
- **Idea:** Churn is not driven by one action, but gradual disengagement patterns

# Window Design and Dataset Construction

- **Backward-sliding windows:**

- Avoid time leakage, capture behavioral dynamics instead of full-sample averages

- **Windows design:**

- Observation window (OBS) of 14 days, recent subwindow (SUB\_OBS) of 7 days
  - Prediction window (PRED) of 10 days
  - Creates 2 backward-sliding snapshots per user

- **Resulting data:**

- Users who churned during observation are omitted since they display churn signals and would not have been flagged in PRED (hidden churn signals – *alive\_users*)
  - User-level data: one observation is one tuple of observation window features and a binary churn flag

# Feature Engineering Philosophy

Activity Volume & Frequency

User Feedback & Interaction

Advertising Exposure

Temporal & Recency Signals

- Features encode changes in behavior, not only levels.
- Zero-imputation when no data is available: NaN songs listened to means zero songs, this is a meaningful signal that should not be dropped.

# Model Comparison

- Models evaluated:
  - Logistic Regression, Random Forest, AdaBoost, LightGBM
- Advantage of boosted tree models:
  - Learns sequentially from errors, captures weak but consistent churn signals
  - Thus, more sample-efficient than Random Forest
- Tradeoff:
  - High flexibility can lead to overfitting
  - Feature selection, regularization, evaluation on a held-out validation set

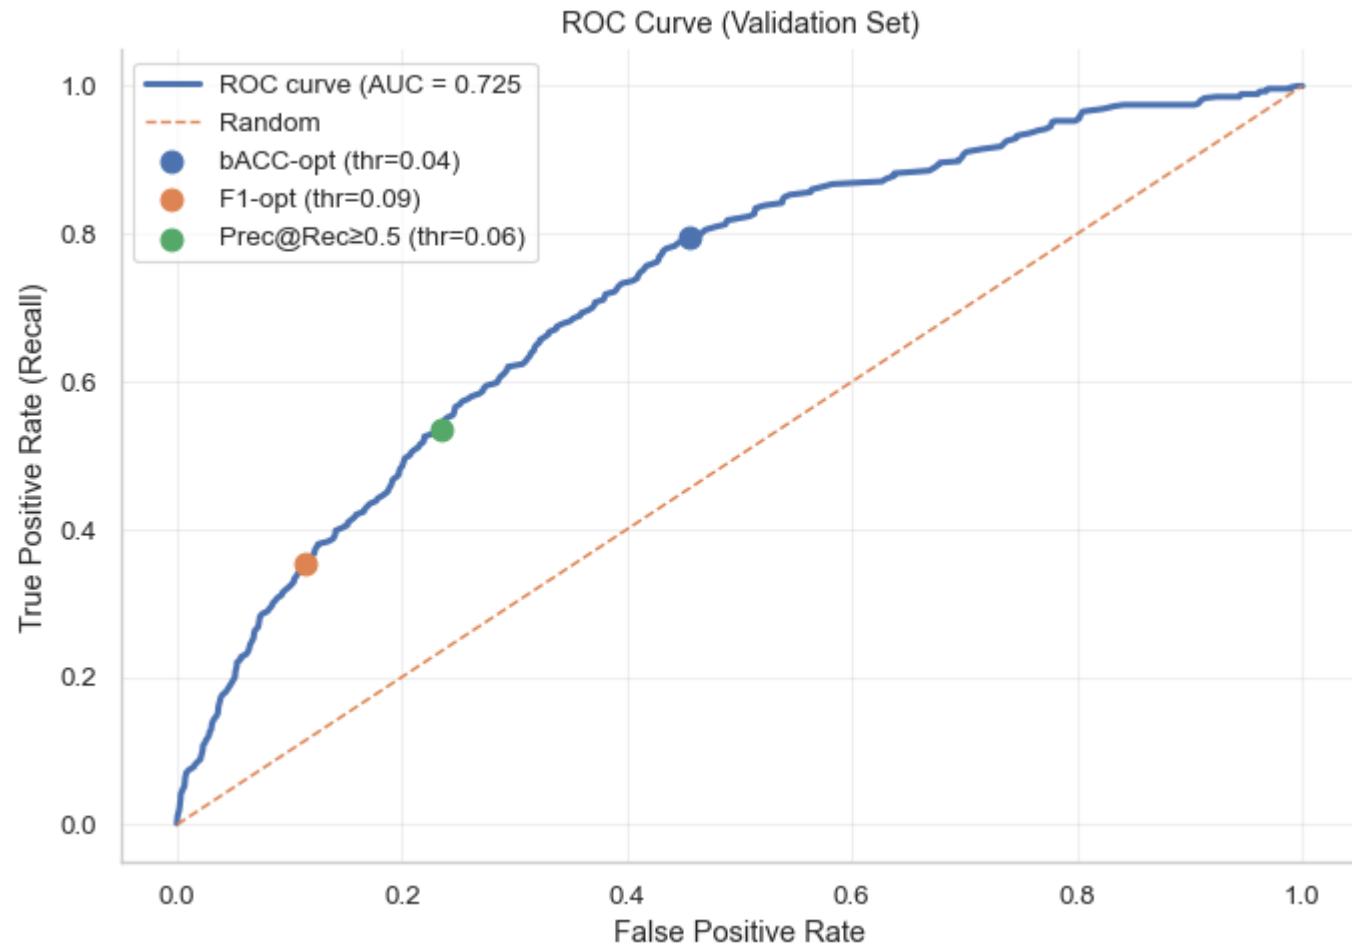
The screenshot shows a machine learning pipeline configuration. The pipeline consists of two main steps: 'SelectKBest' and 'LGBMClassifier'. The 'LGBMClassifier' step is expanded to show its parameters. The parameters are listed in a table with three columns: parameter name, current value, and type. The parameters are:

Parameter	Value	Type
boosting_type	'gbdt'	
num_leaves	31	
max_depth	5	
learning_rate	0.03	
n_estimators	5000	
subsample_for_bin	200000	
objective	'binary'	
class_weight	None	
min_split_gain	0.0	
min_child_weight	0.001	
min_child_samples	100	
subsample	0.8	
subsample_freq	1	
colsample_bytree	0.6	
reg_alpha	0.0	
reg_lambda	10.0	
random_state	42	
n_jobs	-1	
importance_type	'split'	

# Training, Validation , Threshold Tuning

- Training Strategy:
  - Stratified train / validation split (80/20)
  - RandomizedSearchCV with hyperparameter grid and k features
  - SelectKBest: selects k features with highest estimated mutual information with target
    - Reduction in uncertainty about the target after knowing the feature.
- Optimization metric: balanced accuracy
- Best model: LightGBMClassifier to predict validation set churn probabilities
- Threshold Tuning:
  - Default 0.5 is suboptimal for class-imbalanced churn data
  - Tested thresholds: [0.01, 0.95] in steps of 0.015, best threshold was **4%**
  - Balanced accuracy on the held-out validation set with  $\tau = 0.04$  was **67.07%**
  - Threshold tuning is a business-aware decision: do not use raw accuracy, detect potential churners early for higher recall

# Results and Model Behavior



[Threshold tuning by bAcc on held-out validation]				
Best threshold: 0.04				
Best Balanced Accuracy: 0.6707				
	precision	recall	f1-score	support
0	0.98	0.55	0.70	5102
1	0.09	0.80	0.16	275
accuracy			0.56	5377
macro avg	0.53	0.67	0.43	5377
weighted avg	0.93	0.56	0.67	5377

Performance metrics on the held-out validation set

- Test prediction logic:  
Same feature pipeline, use the observation window ending at the cutoff to predict churn in the 10 days after

# Shortcomings of this Model

- Symptoms:
  - Strong validation and public Kaggle performance
  - Weaker private Kaggle score → poor generalization
- Root Causes:
  - Independent feature and hyperparameter selection, iterations do not learn from each other
  - Performance is tied too strongly to validation set and public Kaggle score
- Attempts to fix it:
  - Reduce feature set: without uninformative features
  - There is no single underlying distribution: model users individually, specify user-specific churn models instead of a one-fits-all model
  - Use Optuna study with iterations learning from previous ones
  - Do not rely on the public Kaggle score as proxy for performance on unseen data